

# LLM-Powered Booking Analytics & QA System

## Objective

Develop a system that processes hotel booking data, extracts insights, and enables retrieval-augmented question answering (RAG). The system should provide analytics as mentioned in below sections and answer user queries about the dat.

## Contents

- Dataset Loading
- Data Preprocessing
- Data Analysis and Visaulization
- Processed Data CSV Export

## Links:

- [arindal1 github](#)
- [arindal1 linkedin](#)
- [dataseta](#)

```
!pip install pandas numpy matplotlib seaborn fastapi uvicorn sentence-transformers
transformers faiss-cpu langchain
```

## Data Collection and Pre

```
import pandas as pd
import numpy as np
```

```
# Load the dataset
df = pd.read_csv("data/hotel_bookings.csv")

# Display basic info
print("Dataset Shape:", df.shape)
print(df.info())
print(df.head())
```

```
Dataset Shape: (119390, 32)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 119390 entries, 0 to 119389
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
#   ...
```

---	-----	-----	-----
0	hotel	119390 non-null	object
1	is_canceled	119390 non-null	int64
2	lead_time	119390 non-null	int64
3	arrival_date_year	119390 non-null	int64
4	arrival_date_month	119390 non-null	object
5	arrival_date_week_number	119390 non-null	int64
6	arrival_date_day_of_month	119390 non-null	int64
7	stays_in_weekend_nights	119390 non-null	int64
8	stays_in_week_nights	119390 non-null	int64
9	adults	119390 non-null	int64
10	children	119386 non-null	float64
11	babies	119390 non-null	int64
12	meal	119390 non-null	object
13	country	118902 non-null	object
14	market_segment	119390 non-null	object
15	distribution_channel	119390 non-null	object
16	is_repeated_guest	119390 non-null	int64
17	previous_cancellations	119390 non-null	int64
18	previous_bookings_not_canceled	119390 non-null	int64
19	reserved_room_type	119390 non-null	object
20	assigned_room_type	119390 non-null	object
21	booking_changes	119390 non-null	int64
22	deposit_type	119390 non-null	object
23	agent	103050 non-null	float64
24	company	6797 non-null	float64
25	days_in_waiting_list	119390 non-null	int64
26	customer_type	119390 non-null	object
27	adr	119390 non-null	float64
28	required_car_parking_spaces	119390 non-null	int64
29	total_of_special_requests	119390 non-null	int64
30	reservation_status	119390 non-null	object
31	reservation_status_date	119390 non-null	object

dtypes: float64(4), int64(16), object(12)

memory usage: 29.1+ MB

None

	hotel	is_canceled	lead_time	arrival_date_year	arrival_date_month	\
0	Resort Hotel	0	342	2015	July	
1	Resort Hotel	0	737	2015	July	
2	Resort Hotel	0	7	2015	July	
3	Resort Hotel	0	13	2015	July	
4	Resort Hotel	0	14	2015	July	

	arrival_date_week_number	arrival_date_day_of_month	\
0	27	1	
1	27	1	
2	27	1	
3	27	1	
4	27	1	

stays\_in\_weekend\_nights   stays\_in\_week\_nights   adults   ...   deposit\_type   \

0	0	0	2	...	No Deposit
1	0	0	2	...	No Deposit
2	0	1	1	...	No Deposit
3	0	1	1	...	No Deposit
4	0	2	2	...	No Deposit

	agent	company	days_in_waiting_list	customer_type	adr	\
0	NaN	NaN	0	Transient	0.0	
1	NaN	NaN	0	Transient	0.0	
2	NaN	NaN	0	Transient	75.0	
3	304.0	NaN	0	Transient	75.0	
4	240.0	NaN	0	Transient	98.0	

	required_car_parking_spaces	total_of_special_requests	reservation_status	\
0	0	0	Check-Out	
1	0	0	Check-Out	
2	0	0	Check-Out	
3	0	0	Check-Out	
4	0	1	Check-Out	

	reservation_status_date
0	01-07-15
1	01-07-15
2	02-07-15
3	02-07-15
4	03-07-15

[5 rows x 32 columns]

```
# 1. Convert arrival date columns into a single datetime column.
# Note: 'arrival_date_month' is given as full month name (e.g., "July").
```

```
df['arrival_date'] = pd.to_datetime(
    df['arrival_date_year'].astype(str) + '-' +
    df['arrival_date_month'] + '-' +
    df['arrival_date_day_of_month'].astype(str),
    format='%Y-%B-%d',
    errors='coerce' # in case some dates cannot be parsed
)
```

```
# 2. Convert reservation_status_date into datetime (assumed format: dd-mm-yy)
df['reservation_status_date'] = pd.to_datetime(df['reservation_status_date'],
    format='%d-%m-%y', errors='coerce')
```

```
# 3. Handle missing values:
# - For numerical columns like 'children', fill with 0 or median if
# appropriate.
df['children'] = df['children'].fillna(0)
# - For 'agent' and 'company', fill missing values with a placeholder (e.g., 0
# or "Unknown")
df['agent'] = df['agent'].fillna(0)
df['company'] = df['company'].fillna(0)
# - For 'country', fill missing with "Unknown"
df['country'] = df['country'].fillna("Unknown")
```

```
# 4. Derive additional columns:
# Total number of nights stayed (week + weekend)
df['total_stays'] = df['stays_in_week_nights'] + df['stays_in_weekend_nights']
```

```
# Calculate estimated revenue:
# Revenue = adr * total nights stayed.
# (Assuming each booking's revenue is approximated by the average daily rate times
# total nights.)
df['revenue'] = df['adr'] * df['total_stays']
```

```
# Display the updated dataframe structure
print("\nPreprocessed Data:")
print(df[['arrival_date', 'reservation_status_date', 'total_stays',
'revenue']].head())
```

Preprocessed Data:

	arrival_date	reservation_status_date	total_stays	revenue
0	2015-07-01	2015-07-01	0	0.0
1	2015-07-01	2015-07-01	0	0.0
2	2015-07-01	2015-07-02	1	75.0
3	2015-07-01	2015-07-02	1	75.0
4	2015-07-01	2015-07-03	2	196.0

```
# Function to determine season based on the month (assuming northern hemisphere)
def get_season(date):
    month = date.month
    if month in [12, 1, 2]:
        return 'Winter'
    elif month in [3, 4, 5]:
        return 'Spring'
```

```
elif month in [6, 7, 8]:
    return 'Summer'
else:
    return 'Autumn'

# Apply function to create a new 'season' column (ensure arrival_date is in
datetime format)
df['season'] = df['arrival_date'].apply(lambda x: get_season(x) if pd.notnull(x)
else 'Unknown')

# Verify the new feature
print("Season distribution:")
print(df['season'].value_counts())
```


```
Season distribution:
season
Summer    37477
Spring    32674
Autumn     28462
Winter     20777
Name: count, dtype: int64
```

## Analytics and Visualization

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Group by arrival_date (monthly)
revenue_trend = df.groupby(pd.Grouper(key='arrival_date', freq='M'))
['revenue'].sum()

plt.figure(figsize=(12, 6))
plt.plot(revenue_trend.index, revenue_trend.values, marker='o')
plt.title('Monthly Revenue Trend')
plt.xlabel('Month')
plt.ylabel('Total Revenue')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
# is_canceled: 1 means canceled, 0 means not canceled.
total_bookings = len(df)
total_cancellations = df['is_canceled'].sum()
cancellation_rate = (total_cancellations / total_bookings) * 100

print(f"Overall Cancellation Rate: {cancellation_rate:.2f}%")
```

Overall Cancellation Rate: 37.04%

```
# Count bookings per country
country_counts = df['country'].value_counts().sort_values(ascending=False)


plt.figure(figsize=(12, 6))
sns.barplot(x=country_counts.index[:10], y=country_counts.values[:10],
palette='viridis')
plt.title('Top 10 Countries by Booking Count')
plt.xlabel('Country')
plt.ylabel('Number of Bookings')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

C:\Users\Arindal Char\AppData\Local\Temp\ipykernel\_15688\2707114378.py:5:  
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in

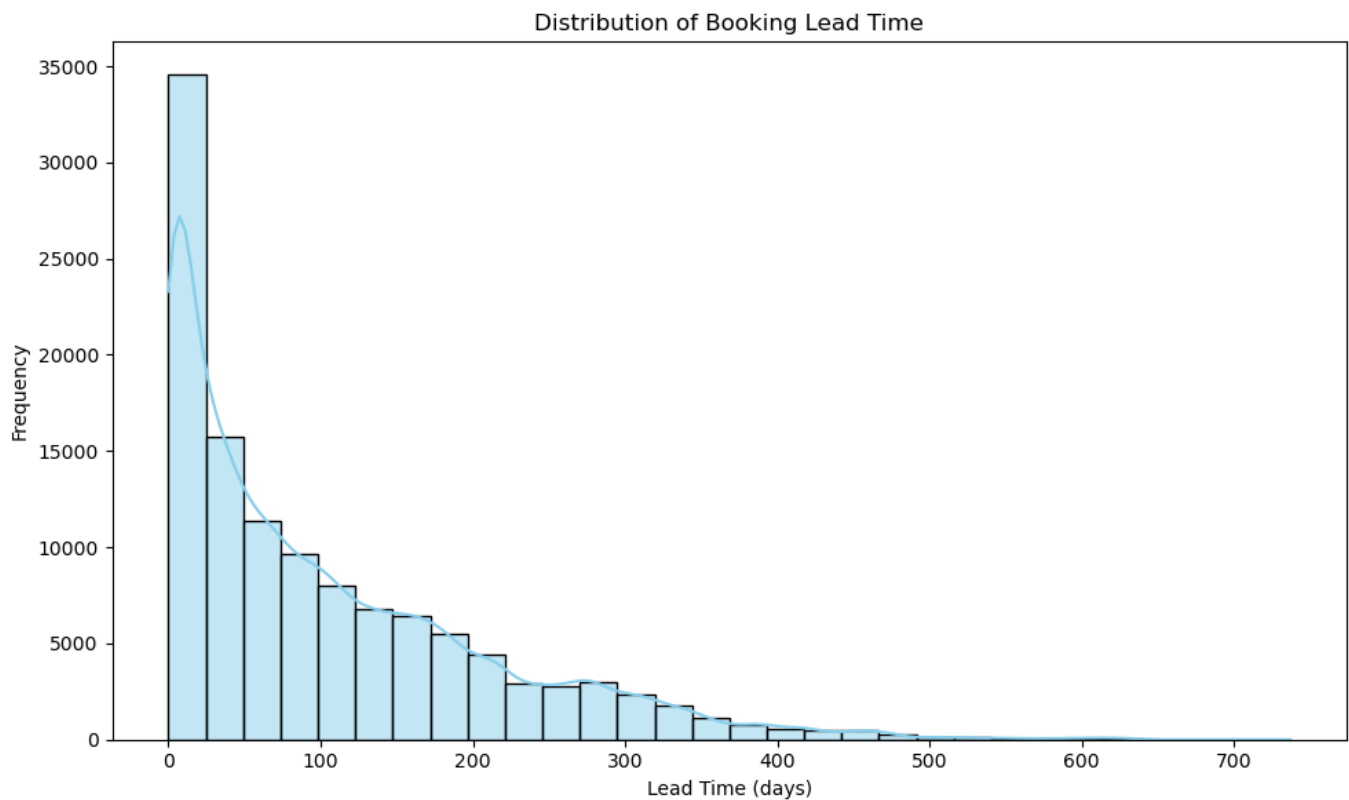
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=country_counts.index[:10], y=country_counts.values[:10],  
palette='viridis')
```



# Histogram of the lead times.

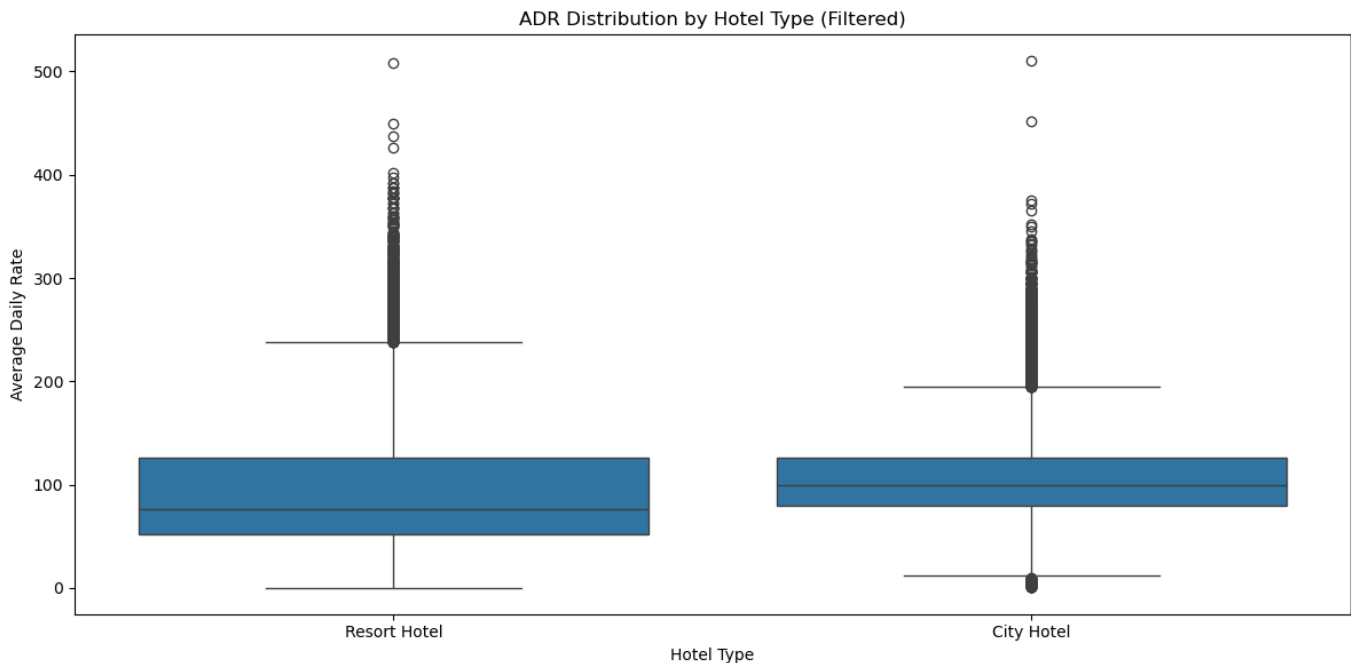
```
plt.figure(figsize=(10, 6))  
sns.histplot(df['lead_time'], bins=30, kde=True, color='skyblue')  
plt.title('Distribution of Booking Lead Time')  
plt.xlabel('Lead Time (days)')  
plt.ylabel('Frequency')  
plt.tight_layout()  
plt.show()
```



```
# Cap the ADR to a maximum value (optional, adjust threshold as per domain
knowledge)
adr_cap = 1000
df['adr_capped'] = df['adr'].apply(lambda x: min(x, adr_cap))
```

```
plt.figure(figsize=(12, 6))
sns.boxplot(x='hotel', y='adr', data=df_filtered)
plt.title('ADR Distribution by Hotel Type (Filtered)')
plt.xlabel('Hotel Type')
plt.ylabel('Average Daily Rate')
plt.tight_layout()
plt.show()
```





```
# Average Revenue per Booking by Customer Type
```

```
# Group by customer_type and calculate average revenue
```

```
avg_revenue_by_customer = df.groupby('customer_type')
```

```
['revenue'].mean().sort_values(ascending=False)
```

```
print("Average Revenue by Customer Type:")
```

```
print(avg_revenue_by_customer)
```

```
plt.figure(figsize=(10, 6))
```

```
sns.barplot(x=avg_revenue_by_customer.index, y=avg_revenue_by_customer.values,
palette='magma')
```

```
plt.title('Average Revenue per Booking by Customer Type')
```

```
plt.xlabel('Customer Type')
```

```
plt.ylabel('Average Revenue')
```

```
plt.tight_layout()
```

```
plt.show()
```

Average Revenue by Customer Type:

customer\_type

Contract 451.196700

Transient 381.631057

Transient-Party 260.493376

Group 243.813328

Name: revenue, dtype: float64

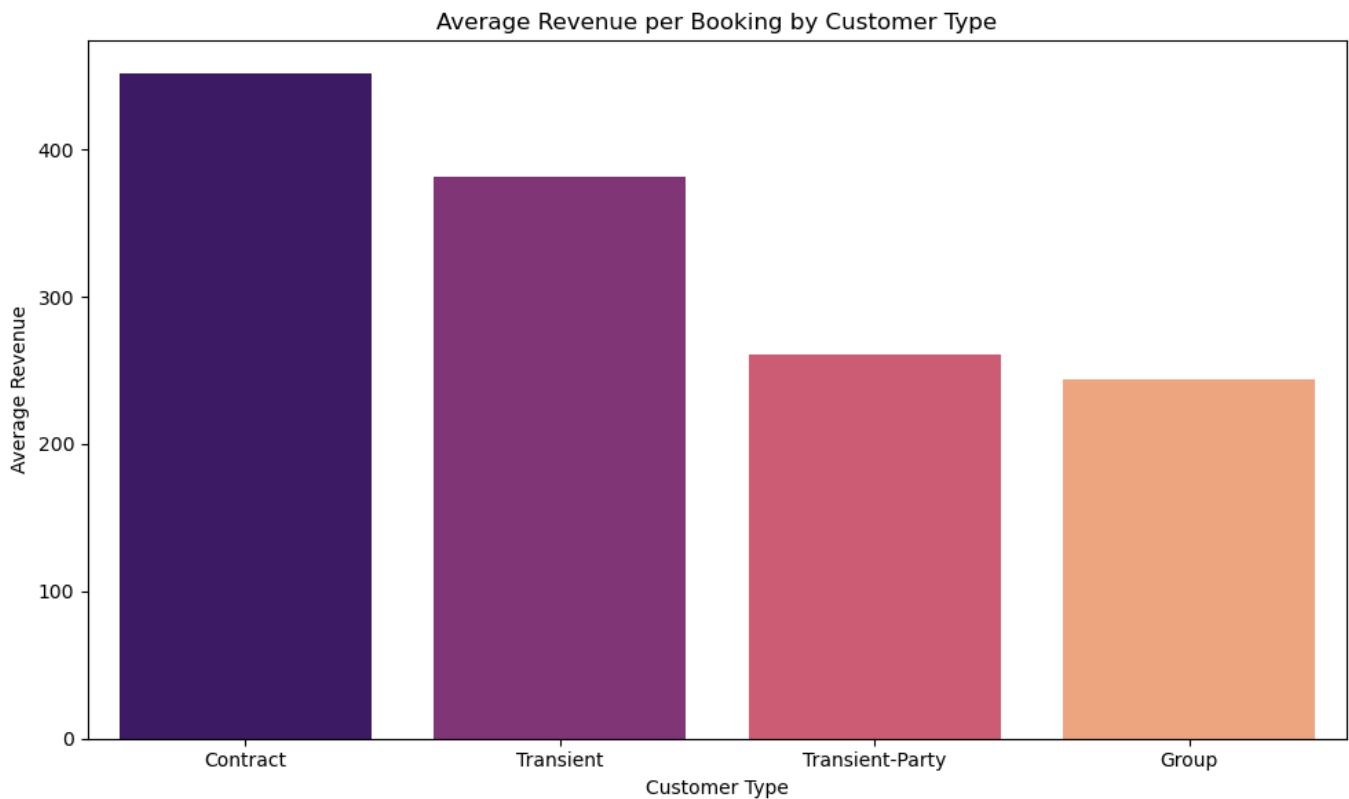
C:\Users\Arindal Char\AppData\Local\Temp\ipykernel\_15688\282020276.py:9:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same

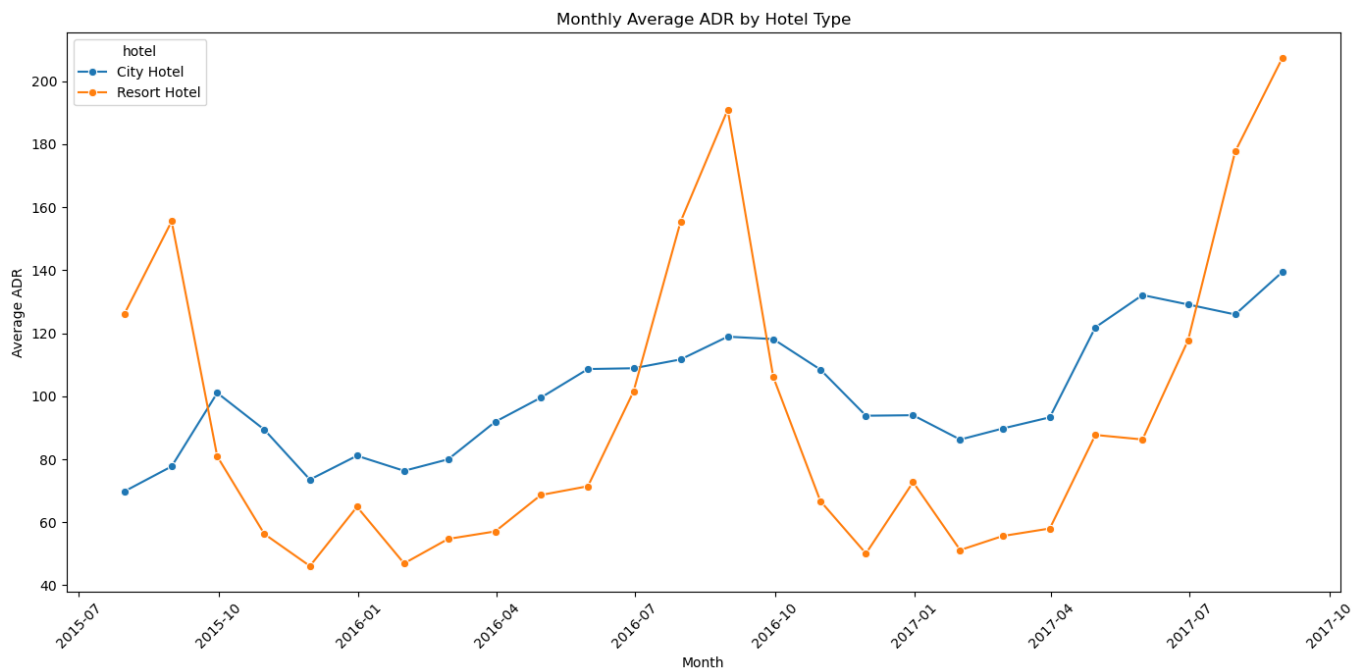
effect.

```
sns.barplot(x=avg_revenue_by_customer.index, y=avg_revenue_by_customer.values,
palette='magma')
```




```
# Group by month and hotel type to calculate mean ADR
adr_trend_hotel = df.groupby([pd.Grouper(key='arrival_date', freq='M'), 'hotel'])
['adr'].mean().reset_index()

plt.figure(figsize=(14,7))
sns.lineplot(data=adr_trend_hotel, x='arrival_date', y='adr', hue='hotel',
marker='o')
plt.title('Monthly Average ADR by Hotel Type')
plt.xlabel('Month')
plt.ylabel('Average ADR')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
# Correlation Heatmap of Numerical Features
numeric_cols = df.select_dtypes(include=[np.number])

plt.figure(figsize=(12, 10))
corr = numeric_cols.corr()
sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm")
plt.title('Correlation Heatmap of Numerical Features')
plt.show()
```



```
# Cancellation Rate by Month and Hotel Type
```

```
# Calculate monthly cancellation rate per hotel type (is_canceled is 1 for canceled, 0 for not canceled)
```

```
cancellation_by_month_hotel = df.groupby([pd.Grouper(key='arrival_date', freq='M'), 'hotel'])['is_canceled'].mean().reset_index()
```

```
plt.figure(figsize=(14,7))
```

```
sns.lineplot(data=cancellation_by_month_hotel, x='arrival_date', y='is_canceled', hue='hotel', marker='o')
```

```
plt.title('Monthly Cancellation Rate by Hotel Type')
```


```
plt.xlabel('Month')
```

```
plt.ylabel('Cancellation Rate')
```

```
plt.xticks(rotation=45)
```

```
plt.tight_layout()
```

```
plt.show()
```



```
# Average Revenue by Season
```

```
# Group by the new 'season' column and calculate average revenue
```

```
avg_revenue_by_season = df.groupby('season')
['revenue'].mean().sort_values(ascending=False)
```

```
print("Average Revenue by Season:")
```

```
print(avg_revenue_by_season)
```

```
plt.figure(figsize=(8,6))
```

```
sns.barplot(x=avg_revenue_by_season.index, y=avg_revenue_by_season.values,
palette='viridis')
```

```
plt.title('Average Revenue by Season')
```

```
plt.xlabel('Season')
```

```
plt.ylabel('Average Revenue')
```

```
plt.tight_layout()
```

```
plt.show()
```

Average Revenue by Season:

season

Summer 507.826623

Spring 320.065285

Autumn 293.079306

Winter 235.464351

Name: revenue, dtype: float64

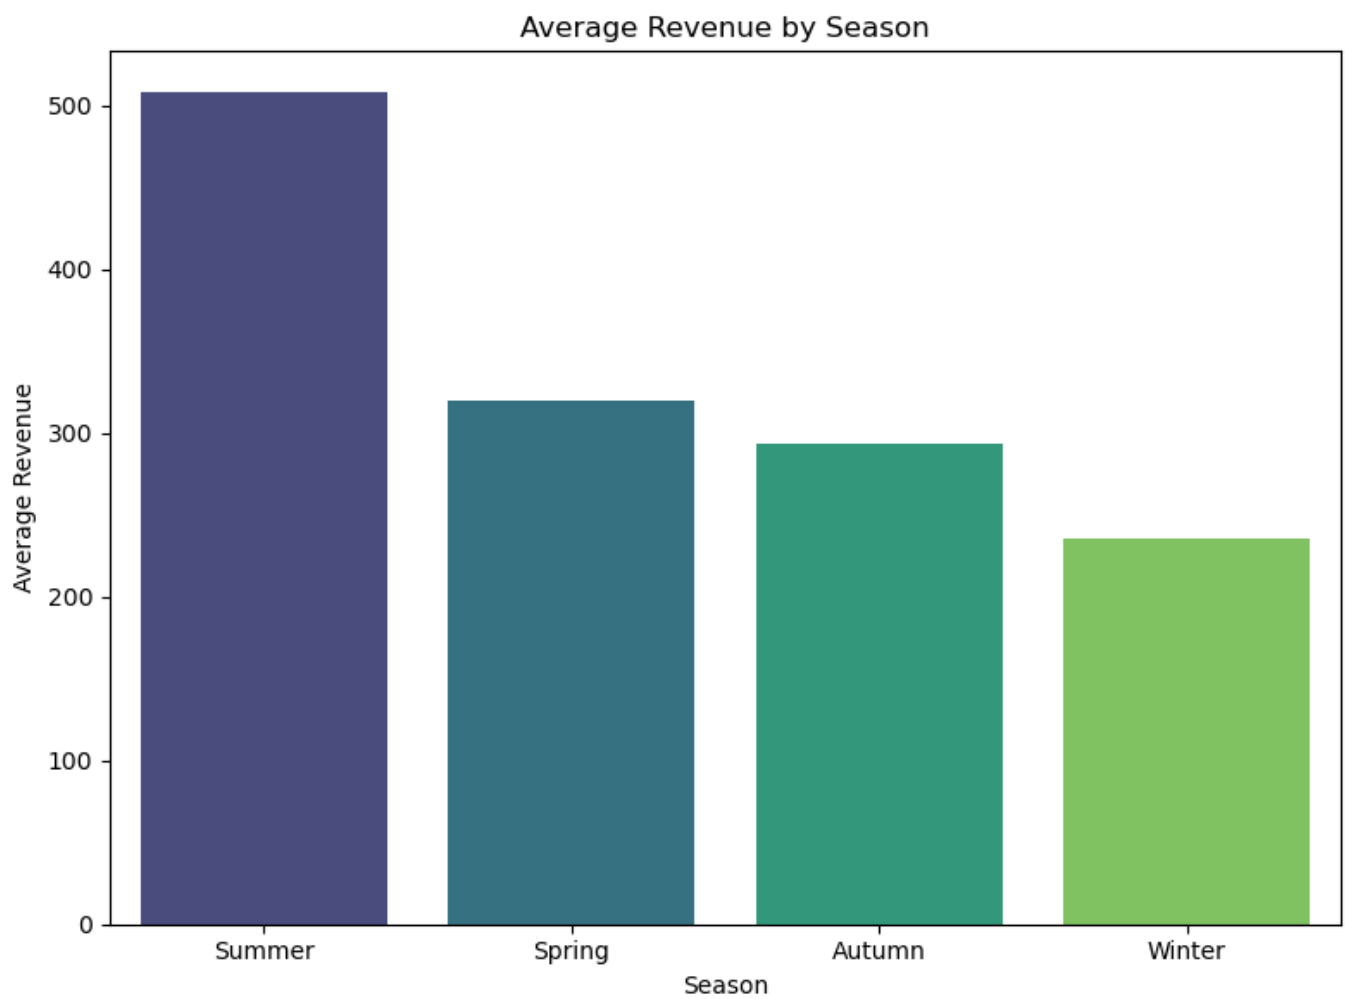
C:\Users\Arindal Char\AppData\Local\Temp\ipykernel\_15688\586029692.py:9:

FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same

effect.

```
sns.barplot(x=avg_revenue_by_season.index, y=avg_revenue_by_season.values,  
palette='viridis')
```



---

## Export Processed Data

```
df.to_csv("hotel_bookings_preprocessed.csv", index=False)
```